

4-FINAL

October 9, 2018

1 Instructions

The following Cells need to be executed.

They are used to download and generate a dataset that has an aggregated count of bike trips per hundredth of an hour through the 24 hours in a day.

I put all this here instead of providing you the dataset directly, so that you could learn something along the way :)

The assignment is in the last cell.

```
In [25]: !pip3 install seaborn
```

Collecting seaborn

Downloading seaborn-0.7.1.tar.gz (158kB)

100% || 163kB 1.8MB/s ta 0:00:01

Building wheels for collected packages: seaborn

Running setup.py bdist_wheel for seaborn ... done

Stored in directory: /Users/mahmoud/Library/Caches/pip/wheels/cb/c8/67/83d615c0ef9b529558525

Successfully built seaborn

Installing collected packages: seaborn

Successfully installed seaborn-0.7.1

1.1 This cell automatically downloads Capital Bikeshare data

```
In [9]: import sys
        sys.path.append('.')
        from utils.bikeshare import download_bikeshare_data

        download_bikeshare_data(2016, 1, '../data/')
```

Downloading: 2016 Q1 | Extracting... | Created: ../data/2016-Q1-cabi-trip-history-data.csv

1.1.1 And here we read in the data

```
In [4]: import seaborn as sns
        import matplotlib.pyplot as plt
        %matplotlib inline
```

```
import pandas as pd
bikes = pd.read_csv('../data/2016Q1-capitalbikeshare-tripdata.csv')
bikes.head()
bikes['start'] = pd.to_datetime(bikes['Start date'], infer_datetime_format=True)
bikes['end'] = pd.to_datetime(bikes['End date'], infer_datetime_format=True)
```

```
bikes.head()
```

```
Out[4]:
```

	Duration	Start date	End date	Start station number	\
0	166	2016-01-01 00:06:58	2016-01-01 00:09:44	31102	
1	448	2016-01-01 00:10:20	2016-01-01 00:17:48	32039	
2	715	2016-01-01 00:13:52	2016-01-01 00:25:48	31222	
3	213	2016-01-01 00:15:29	2016-01-01 00:19:03	31506	
4	872	2016-01-01 00:16:16	2016-01-01 00:30:49	31041	

	Start station	End station number	\
0	11th & Kenyon St NW	31105	
1	Old Georgetown Rd & Southwick St	32002	
2	New York Ave & 15th St NW	31214	
3	1st & Rhode Island Ave NW	31509	
4	Prince St & Union St	31048	

	End station	Bike number	Member type	start	\
0	14th & Harvard St NW	W01346	Member	2016-01-01 00:06:58	
1	Bethesda Ave & Arlington Rd	W22202	Member	2016-01-01 00:10:20	
2	17th & Corcoran St NW	W21427	Member	2016-01-01 00:13:52	
3	New Jersey Ave & R St NW	W01294	Member	2016-01-01 00:15:29	
4	King St Metro South	W22058	Member	2016-01-01 00:16:16	

	end
0	2016-01-01 00:09:44
1	2016-01-01 00:17:48
2	2016-01-01 00:25:48
3	2016-01-01 00:19:03
4	2016-01-01 00:30:49

1.1.2 Create a new column that represents the hour of the day

```
In [5]: bikes['hour_of_day'] = (bikes.start.dt.hour + (bikes.start.dt.minute/60).round(2))
```


1.1.3 Aggregate to get a count per hour/minute of the day across all trips

```
In [6]: hours = bikes.groupby('hour_of_day').agg('count')
hours['hour'] = hours.index
```

```
hours.start.plot()
# import seaborn as sns
```

```
# sns.lmplot(x='hour', y='start', data=hours, aspect=1.5, scatter_kws={'alpha':0.2})
```

Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x27b0f91c208>



2 Assignment 4

Using the `hours` dataframe and the `hour_of_day` column, perform the following cells.

Explain the results in a **paragraph + charts** of to describe which model you'd recommend

2.1 1. Create 3 models fit to `hour_of_day` with varying polynomial degrees

```
In [7]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import linear_model
```

```
In [15]: hours
hours.head()
```

```
Out[15]:
```

hour_of_day	Duration	Start date	End date	Start station number	\
0.00	71	71	71	71	
0.02	88	88	88	88	
0.03	74	74	74	74	
0.05	77	77	77	77	
0.07	66	66	66	66	


hour_of_day	Start station	End station number	End station	Bike number \
0.00	71	71	71	71
0.02	88	88	88	88
0.03	74	74	74	74
0.05	77	77	77	77
0.07	66	66	66	66

hour_of_day	Member type	start	end	hour
0.00	71	71	71	0.00
0.02	88	88	88	0.02
0.03	74	74	74	0.03
0.05	77	77	77	0.05
0.07	66	66	66	0.07

```
In [19]: n = 552400
x = np.linspace(0.01, 3, n).reshape(-1, 1)
y = np.linspace(0.01, 3, n) * np.linspace(0.01, 3, n) + np.random.rand(n) - .5
```

```
plt.scatter(x,y)
plt.plot(x, x*linear.coef_ + linear.intercept_, c='b')
```

```
Out[19]: [<matplotlib.lines.Line2D at 0x27b0fafa9b0>]
```



```
In [17]: linear = linear_model.LinearRegression()
```

```
linear.fit(x, y)
```


```
linear.coef_, linear.intercept_
```

```
Out[17]: (array([3.00979324]), -1.5202046888365257)
```

```
In [18]: plt.scatter(x,y)
```

```
plt.plot(x, x*linear.coef_ + linear.intercept_, c='b')
```

```
Out[18]: [<matplotlib.lines.Line2D at 0x27b0fb5d3c8>]
```



```
In [21]: from sklearn.preprocessing import PolynomialFeatures
```

```
poly = PolynomialFeatures(degree=15)
```

```
x_15= poly.fit_transform(x.reshape(-1, 1))
```

```
In [22]: linear = linear_model.LinearRegression()
```

```
linear.fit(x_15, y)
```

```
(linear.coef_, linear.intercept_)
```

```
Out[22]: (array([ 0.00000000e+00,  6.65557172e-01, -9.74062628e+00,  7.81070563e+01,  
                 -3.18460897e+02,  8.23299881e+02, -1.45031105e+03,  1.81566889e+03,
```

```

-1.65301894e+03,  1.10494527e+03, -5.41441591e+02,  1.91966499e+02,
-4.78655265e+01,  7.95204148e+00, -7.89550837e-01,  3.54187776e-02]),
-0.01200074100206594)

```

```


In [23]: plt.scatter(x,y)
plt.plot(x, np.dot(x_15, linear.coef_) + linear.intercept_, c='b')

```

```

Out[23]: [<matplotlib.lines.Line2D at 0x27b10a8a6d8>]

```



```

In [24]: from sklearn.preprocessing import PolynomialFeatures

```

```

poly = PolynomialFeatures(degree=2)

x_2= poly.fit_transform(x.reshape(-1, 1))

```

```

In [25]: linear = linear_model.LinearRegression()

```

```

linear.fit(x_2, y)

(linear.coef_, linear.intercept_)

```

```

Out[25]: (array([ 0.00000000e+00, -2.80547562e-04,  1.00001556e+00]),
 0.0002031101826003301)


```

```

In [26]: plt.scatter(x,y)
plt.plot(x, np.dot(x_2, linear.coef_) + linear.intercept_, c='b')

```

Out [26]: [



2.2 2. Choose one of the polynomial models and create 3 new models fit to hour_of_day with different Ridge Regression α (alpha) Ridge Coefficient values

In [28]: `from sklearn.preprocessing import PolynomialFeatures`

```
poly = PolynomialFeatures(degree=2)

x_2= poly.fit_transform(x.reshape(-1, 1))
```

In [29]: `linear = linear_model.LinearRegression()`

```
linear.fit(x_2, y)

(linear.coef_, linear.intercept_)
```

Out [29]: `(array([0.00000000e+00, -2.80547562e-04, 1.00001556e+00]),
 0.0002031101826003301)`

In [30]: `ridge = linear_model.Ridge()`


```
ridge.fit(x_2, y)

(ridge.coef_, ridge.intercept_)
```

```
Out [30]: (array([ 0.00000000e+00, -2.68265367e-04,  1.00001148e+00]),
          0.00019690716894382732)
```

```
In [31]: plt.scatter(x,y)
plt.plot(x, np.dot(x_2, linear.coef_) + linear.intercept_, c='b')
plt.plot(x, np.dot(x_2, ridge.coef_) + ridge.intercept_, c='r')
```

```
Out [31]: [<matplotlib.lines.Line2D at 0x27b21c3de48>]
```



```
In [33]: from sklearn.preprocessing import PolynomialFeatures
```

```
poly = PolynomialFeatures(degree=10)
```

```
x_10= poly.fit_transform(x.reshape(-1, 1))
```

```
In [34]: linear = linear_model.LinearRegression()
```

```
linear.fit(x_10, y)
```

```
(linear.coef_, linear.intercept_)
```

```
Out [34]: (array([ 0.00000000e+00, -2.69350137e-01,  3.23020591e+00, -7.93397599e+00,
                  1.50586040e+01, -1.69134968e+01,  1.17999450e+01, -5.16990652e+00,
                  1.38409188e+00, -2.06961823e-01,  1.32467576e-02]),
          0.00821148524487425)
```



```
In [35]: ridge = linear_model.Ridge()
```

```
ridge.fit(x_10, y)
```

```
(ridge.coef_, ridge.intercept_)
```


```
Out[35]: (array([ 0.00000000e+00,  5.75627166e-02,  8.15267892e-01,  2.14350609e-01,  
                -5.16450114e-02, -6.66603686e-02,  2.71026942e-02,  2.37241422e-02,  
                -2.01040579e-02,  5.41370221e-03, -5.11622599e-04]),  
         -0.004444042052149074)
```

```
In [36]: plt.scatter(x,y)
```

```
plt.plot(x, np.dot(x_10, linear.coef_) + linear.intercept_, c='b')
```

```
plt.plot(x, np.dot(x_10, ridge.coef_) + ridge.intercept_, c='r')
```

```
Out[36]: [<matplotlib.lines.Line2D at 0x27b21cc32e8>]
```



```
In [41]: from sklearn.preprocessing import PolynomialFeatures
```

```
poly = PolynomialFeatures(degree=50)
```

```
x_50= poly.fit_transform(x.reshape(-1, 1))
```

```
In [43]: linear = linear_model.LinearRegression()
```

```
linear.fit(x_50, y)
```

```
(linear.coef_, linear.intercept_)
```

```
Out [43]: (array([-7.84092274e-15, -6.28289279e-09,  3.25145576e-10, -1.57611314e-11,
-1.00807983e-10,  4.96783680e-11, -5.49029186e-11,  2.50288708e-11,
-1.35409544e-11, -8.94452050e-11, -3.41023130e-10, -2.63362808e-10,
 2.49462510e-13,  4.93069187e-13,  9.62038249e-13,  1.85377499e-12,
 3.52796551e-12,  6.62948500e-12,  1.22939988e-11,  2.24816467e-11,
 4.04989290e-11,  7.17774222e-11,  1.24964481e-10,  2.13315392e-10,
 3.56208461e-10,  5.80264770e-10,  9.18973083e-10,  1.40888850e-09,
 2.07957683e-09,  2.93412479e-09,  3.91854130e-09,  4.88384816e-09,
 5.55633790e-09,  5.54926479e-09,  4.46543565e-09,  2.13238582e-09,
-1.05868097e-09, -3.90481767e-09, -4.64164904e-09, -2.10766202e-09,
 2.43784701e-09,  4.71228663e-09,  1.01570546e-09, -4.68704963e-09,
-1.71789488e-09,  5.96949340e-09, -4.25118396e-09,  1.54375889e-09,
-3.18034208e-10,  3.55562964e-11, -1.68615021e-12]),
1.3545903246543278)
```

```
In [44]: ridge = linear_model.Ridge()
```

```
ridge.fit(x_50, y)
```

```
(ridge.coef_, ridge.intercept_)
```

```
Out [44]: (array([-4.54651451e-10, -5.45183013e-08,  1.45916922e-08,  3.79095648e-08,
-1.03128319e-08,  4.94298116e-09, -6.07296967e-10, -1.24841405e-09,
 1.72398705e-09, -4.95932816e-10,  4.62597221e-08, -4.86216542e-08,
 2.95555274e-09, -3.39590468e-09, -5.39262619e-09,  9.22836668e-10,
 5.96188294e-09,  9.43570327e-11,  1.57776112e-10,  2.84270233e-10,
 4.27793062e-10,  4.93504949e-10,  8.14715766e-10,  1.16570418e-09,
 1.57830235e-09,  2.14680423e-09,  2.94137260e-09,  3.61328913e-09,
 4.28970383e-09,  4.76800546e-09,  4.55545859e-09,  3.77685886e-09,
 2.31607521e-09, -2.75192354e-12, -2.26865886e-09, -3.63123481e-09,
-3.05716080e-09, -4.37109221e-10,  2.76484288e-09,  3.50522401e-09,
 2.77570365e-10, -3.32272907e-09, -2.30737298e-09,  3.82373934e-09,
 5.26011859e-10, -2.92605503e-09,  2.04301439e-09, -7.22374677e-10,
 1.45758086e-10, -1.60650606e-11,  7.55336314e-13]), 2.583655266060578)
```

```
In [45]: plt.scatter(x,y)
```

```
plt.plot(x, np.dot(x_50, linear.coef_) + linear.intercept_, c='b')
```

```
plt.plot(x, np.dot(x_50, ridge.coef_) + ridge.intercept_, c='r')
```

```
Out [45]: [<matplotlib.lines.Line2D at 0x27b21d24cf8>]
```

