

12-FINAL

December 18, 2018

1 Assignment 12 - Neural Networks image recognition

Use both MLNN and the ConvNet to solve the following problem.

1. Add random noise (i.e. `np.random.normal`) to the images in training and testing. Make sure each image gets a different noise feature added to it. Inspect by printing out an image.
2. Compare the loss/accuracy (train, val) after N epochs for both MLNN and ConvNet with and without noise.
3. Vary the amount of noise (multiply `np.random.normal` by a factor) and keep track of the accuracy and loss (for training and validation) and plot these results.

2 Neural Networks - Image Recognition

```
In [1]: import keras
        from keras.datasets import mnist
        from keras.models import Sequential
        from keras.optimizers import RMSprop
        from keras.layers import Dense, Dropout, Flatten
        from keras.layers import Conv2D, MaxPooling2D
        from keras import backend
```

```
C:\Users\Erin\Anaconda3\lib\site-packages\h5py\__init__.py:36: FutureWarning: Conversion of the
  from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
In [10]: import matplotlib.pyplot as plt
         %matplotlib inline
```


2.1 Multi Layer Neural Network

Trains a simple deep NN on the MNIST dataset. Gets to 98.40% test accuracy after 20 epochs (there is *a lot* of margin for parameter tuning).

```
In [11]: # the data, shuffled and split between train and test sets
         (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [18]: plt.imshow(x_train[2])
```

Out[18]: <matplotlib.image.AxesImage at 0x2174a9b35f8>



```
In [19]: import numpy as np
         np.max(x_train[2])
```

Out[19]: 255

```
In [16]: y_train[2]
```

Out[16]: 4

```
In [24]: # the data, shuffled and split between train and test sets
         (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
         x_train = x_train.reshape(60000, 784)
         x_test = x_test.reshape(10000, 784)
         x_train = x_train.astype('float32')
         x_test = x_test.astype('float32')
         x_train /= 255
         x_test /= 255
         print(x_train.shape[0], 'train samples')
         print(x_test.shape[0], 'test samples')
```

60000 train samples

10000 test samples

[illegible]

0. , 0.07058824, 0.85882354, 0.99215686, 0.99215686,
 0.99215686, 0.99215686, 0.99215686, 0.7764706 , 0.7137255 ,
 0.96862745, 0.94509804, 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0.3137255 , 0.6117647 , 0.41960785, 0.99215686, 0.99215686,
 0.8039216 , 0.04313726, 0. , 0.16862746, 0.6039216 ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.05490196,
 0.00392157, 0.6039216 , 0.99215686, 0.3529412 , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.54509807,
 0.99215686, 0.74509805, 0.00784314, 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0.04313726, 0.74509805, 0.99215686,
 0.27450982, 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.13725491, 0.94509804, 0.88235295, 0.627451 ,
 0.42352942, 0.00392157, 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0.31764707, 0.9411765 , 0.99215686, 0.99215686, 0.46666667,
 0.09803922, 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.1764706 ,
 0.7294118 , 0.99215686, 0.99215686, 0.5882353 , 0.10588235,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0.0627451 , 0.3647059 ,
 0.9882353 , 0.99215686, 0.73333335, 0. , 0. ,

0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.9764706	, 0.99215686,	
0.9764706	, 0.2509804	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.18039216,	0.50980395,	
0.7176471	, 0.99215686,	0.99215686,	0.8117647	, 0.00784314,	
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.15294118,	
0.5803922	, 0.8980392	, 0.99215686,	0.99215686,	0.99215686,	
0.98039216,	0.7137255	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.09411765,	0.44705883,	0.8666667	, 0.99215686,	0.99215686,	
0.99215686,	0.99215686,	0.7882353	, 0.30588236,	0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.09019608,	0.25882354,	0.8352941	, 0.99215686,	
0.99215686,	0.99215686,	0.99215686,	0.7764706	, 0.31764707,	
0.00784314,	0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.07058824,	0.67058825,	0.85882354,	
0.99215686,	0.99215686,	0.99215686,	0.99215686,	0.7647059	,
0.3137255	, 0.03529412,	0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.21568628,	0.6745098	,
0.8862745	, 0.99215686,	0.99215686,	0.99215686,	0.99215686,	
0.95686275,	0.52156866,	0.04313726,	0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.53333336,	0.99215686,	0.99215686,	0.99215686,	
0.83137256,	0.5294118	, 0.5176471	, 0.0627451	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,

[illegible]

```
In [22]: keras.utils.to_categorical(y_train, num_classes)
```

```
Out[22]: array([[1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 ...,
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.]],

                [[0., 1., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 ...,
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.]],

                [[1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 ...,
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.]],

                ...],

                [[1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 ...,
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.],
                 [1., 0., 0., ..., 0., 0., 0.]])
```

```

[1., 0., 0., ..., 0., 0., 0.],
[1., 0., 0., ..., 0., 0., 0.],
[1., 0., 0., ..., 0., 0., 0.]],

[[1., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.],
 ...,
 [1., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.]],

[[1., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.],
 ...,
 [1., 0., 0., ..., 0., 0., 0.],
 [0., 1., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.]]], dtype=float32)

```

```
In [23]: y_train
```

```
Out[23]: array([[0., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 1., 0.]], dtype=float32)
```

```
In [25]: batch_size = 128
num_classes = 10
epochs = 20
```

```

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

```

```

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 512)	262656
dropout_4 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 10)	5130

```

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```

```

Train on 60000 samples, validate on 10000 samples

```

```

Epoch 1/20
60000/60000 [=====] - 7s 110us/step - loss: 0.2433 - acc: 0.9250 - val.
Epoch 2/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1024 - acc: 0.9690 - val.
Epoch 3/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0735 - acc: 0.9773 - val.
Epoch 4/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0597 - acc: 0.9823 - val.
Epoch 5/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0498 - acc: 0.9853 - val.
Epoch 6/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0426 - acc: 0.9872 - val.
Epoch 7/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0366 - acc: 0.9888 - val.
Epoch 8/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0334 - acc: 0.9903 - val.
Epoch 9/20

```



```

60000/60000 [=====] - 6s 94us/step - loss: 0.0301 - acc: 0.9913 - val.
Epoch 10/20
60000/60000 [=====] - 5s 92us/step - loss: 0.0304 - acc: 0.9913 - val.
Epoch 11/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0263 - acc: 0.9925 - val.
Epoch 12/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0257 - acc: 0.9928 - val.
Epoch 13/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0240 - acc: 0.9930 - val.
Epoch 14/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0218 - acc: 0.9939 - val.
Epoch 15/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0220 - acc: 0.9941 - val.
Epoch 16/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0189 - acc: 0.9947 - val.
Epoch 17/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0197 - acc: 0.9948 - val.
Epoch 18/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0183 - acc: 0.9954 - val.
Epoch 19/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0170 - acc: 0.9954 - v
Epoch 20/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0180 - acc: 0.9958 - val.
Test loss: 0.10086243585582774
Test accuracy: 0.9852

```

2.2 Conv Net

Trains a simple convnet on the MNIST dataset. Gets to 99.25% test accuracy after 12 epochs (there is still a lot of margin for parameter tuning).

```

In [26]: # input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')

```

```

x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

```

```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

```

In [27]: batch_size = 128
        num_classes = 10
        epochs = 12

```

```

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

```

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
        batch_size=batch_size,
        epochs=epochs,
        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 94s 2ms/step - loss: 0.2796 - acc: 0.9139 - val.

```

Epoch 2/12
60000/60000 [=====] - 99s 2ms/step - loss: 0.0875 - acc: 0.9744 - val.
Epoch 3/12
60000/60000 [=====] - 98s 2ms/step - loss: 0.0677 - acc: 0.9798 - val.
Epoch 4/12
60000/60000 [=====] - 97s 2ms/step - loss: 0.0548 - acc: 0.9839 - val.
Epoch 5/12
60000/60000 [=====] - 112s 2ms/step - loss: 0.0469 - acc: 0.9861 - val.
Epoch 6/12
60000/60000 [=====] - 106s 2ms/step - loss: 0.0409 - acc: 0.9878 - val.
Epoch 7/12
60000/60000 [=====] - 90s 1ms/step - loss: 0.0387 - acc: 0.9882 - val.
Epoch 8/12
60000/60000 [=====] - 106s 2ms/step - loss: 0.0348 - acc: 0.9893 - val.
Epoch 9/12
60000/60000 [=====] - 117s 2ms/step - loss: 0.0330 - acc: 0.9898 - val.
Epoch 10/12
60000/60000 [=====] - 86s 1ms/step - loss: 0.0295 - acc: 0.9906 - val.
Epoch 11/12
60000/60000 [=====] - 86s 1ms/step - loss: 0.0299 - acc: 0.9906 - val.
Epoch 12/12
60000/60000 [=====] - 86s 1ms/step - loss: 0.0281 - acc: 0.9914 - val.
Test loss: 0.026129956484261856
Test accuracy: 0.9928

```

```

In [28]: score = model.evaluate(x_test, y_test, verbose=0)
         print('Test loss:' , score[0])
         print('Test accuracy:' , score[1])

```

```

Test loss: 0.026129956484261856
Test accuracy: 0.9928

```

```


In [29]: plt.imshow(x_test[0].reshape(28,28))

```

```

Out[29]: <matplotlib.image.AxesImage at 0x2176ea9ee48>

```



```
In [30]: model.predict(x_test[0:1]), y_test[:1]
```

```
Out[30]: (array([[1.1235774e-12, 3.6403717e-11, 4.6634319e-10, 1.3857475e-09,
                  2.0686898e-11, 3.6422727e-12, 1.8520169e-16, 1.0000000e+00,
                  1.0492592e-11, 5.4979825e-09]], dtype=float32),
          array([[0., 0., 0., 0., 0., 0., 0., 1., 0., 0.]], dtype=float32))
```

```
In [31]: model.predict_classes(x_test[0:1]), y_test[:1]
```

```
Out[31]: (array([7], dtype=int64),
          array([[0., 0., 0., 0., 0., 0., 0., 1., 0., 0.]], dtype=float32))
```

```
In [40]: # CovNet with noise
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
```

```
np.random.normal(0, 0.1, 1)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```